

```
In [1]: 1 import numpy as np
2 import pandas as pd
3
4 import matplotlib.pyplot as plt
5 %matplotlib inline
6
7 import seaborn as sns
8
9 import warnings
10 warnings.filterwarnings('ignore')
```

```
In [2]: 1 pd.set_option('display.max_columns', None)
2 pd.set_option('display.max_rows', None)
```

```
In [3]: 1 cars = pd.read_csv('C:\\Users\\gauta\\Desktop\\AI_ML\\ML2\\Advance Regression\\Advanced-Regression-main\\Regularization\\Reg
2 cars.head()
```

Out[3]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	carheight	curbwei
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	20
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	20
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	52.4	20
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	54.3	20
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	54.3	20

```
In [4]: 1 cars.shape
```

Out[4]: (205, 26)

```
In [5]: 1 cars.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling             205 non-null   int64
2   CarName               205 non-null   object
3   fueltype              205 non-null   object
4   aspiration            205 non-null   object
5   doornumber            205 non-null   object
6   carbody               205 non-null   object
7   drivewheel            205 non-null   object
8   enginelocation        205 non-null   object
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
14  enginetype            205 non-null   object
15  cylindernumber        205 non-null   object
16  enginesize            205 non-null   int64
17  fuelsystem            205 non-null   object
18  boreratio             205 non-null   float64
19  stroke                205 non-null   float64
20  compressionratio      205 non-null   float64
21  horsepower            205 non-null   int64
22  peakrpm               205 non-null   int64
23  citympg               205 non-null   int64
24  highwaympg            205 non-null   int64
25  price                 205 non-null   float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

```
In [6]: 1 cars.isnull().sum().sum()
```

Out[6]: 0

```
In [7]: 1 cars['symboling'].value_counts()
```

```
Out[7]: 0    67
        1    54
        2    32
        3    27
       -1    22
       -2     3
        Name: symboling, dtype: int64
```

```
In [8]: 1 cars.columns
```

```
Out[8]: Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',
              'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',
              'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',
              'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',
              'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',
              'price'],
              dtype='object')
```

```
In [9]: 1 for column in cars.columns:
        2     print(column)
```

```
car_ID
symboling
CarName
fueltype
aspiration
doornumber
carbody
drivewheel
enginelocation
wheelbase
carlength
carwidth
carheight
curbweight
enginetype
cylindernumber
enginesize
fuelsystem
boreratio
stroke
compressionratio
horsepower
peakrpm
citympg
highwaympg
price
```

```
In [10]: 1 for column, value in cars.iteritems():
        2     print(column)
```

```
car_ID
symboling
CarName
fueltype
aspiration
doornumber
carbody
drivewheel
enginelocation
wheelbase
carlength
carwidth
carheight
curbweight
enginetype
cylindernumber
enginesize
fuelsystem
boreratio
stroke
compressionratio
horsepower
peakrpm
citympg
highwaympg
price
```

```
In [11]: 1 for column, value in cars.iteritems():
2         if value.dtypes == 'object':
3             print(column)
```

CarName
fueltype
aspiration
doornumber
carbody
drivewheel
enginelocation
enginetype
cylindernumber
fuelsystem

```
In [12]: 1 cars.select_dtypes(include=object).columns
```

```
Out[12]: Index(['CarName', 'fueltype', 'aspiration', 'doornumber', 'carbody',
               'drivewheel', 'enginelocation', 'enginetype', 'cylindernumber',
               'fuelsystem'],
              dtype='object')
```

```
In [13]: 1 cars.select_dtypes(include=object).columns.tolist()
```

```
Out[13]: ['CarName',
          'fueltype',
          'aspiration',
          'doornumber',
          'carbody',
          'drivewheel',
          'enginelocation',
          'enginetype',
          'cylindernumber',
          'fuelsystem']
```

```
In [14]: 1 cars.select_dtypes(include=object).columns.value_counts()
```

```
Out[14]: CarName      1
fueltype      1
aspiration      1
doornumber      1
carbody      1
drivewheel      1
enginelocation  1
enginetype      1
cylindernumber  1
fuelsystem      1
dtype: int64
```

```
In [15]: 1 for column in cars.select_dtypes(include=object):
2         #     if cars.columns.dtypes == 'object':
3             print(column)
```

CarName
fueltype
aspiration
doornumber
carbody
drivewheel
enginelocation
enginetype
cylindernumber
fuelsystem

```
In [16]: 1 # for i in heart.columns:
2 #     x = heart[i].value_counts()
3 #     print("Column name is:",i,"and it value is:",x)
4
5 for i in cars.select_dtypes(include=object):
6     x = cars[i].value_counts()
7     print("Column name is:",i)
8     print(x)
9     print('-----')
```

```
Column name is: CarName
toyota corona          6
toyota corolla         6
peugeot 504            6
subaru dl              4
mitsubishi mirage g4   3
mazda 626              3
toyota mark ii         3
mitsubishi outlander   3
mitsubishi g4          3
honda civic            3
volvo 264gl            2
bmw 320i               2
isuzu D-Max            2
audi 100ls             2
volvo 244dl            2
porsche cayenne        2
toyota corolla liftback 2
honda accord           2
^
```

```
In [17]: 1 cars.head()
```

Out[17]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	carheight	curbwei
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	21
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	21
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	52.4	21
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	54.3	21
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	54.3	21

```
In [18]: 1 cars.carlength.describe()
```

Out[18]:

```
count    205.000000
mean     174.049268
std       12.337289
min      141.100000
25%      166.300000
50%      173.200000
75%      183.100000
max      208.100000
Name: carlength, dtype: float64
```

```
In [19]: 1 # cars.select_dtypes(include=object).columns.tolist()
```

```
In [20]: 1 for i in cars.select_dtypes(include=object):
2         print(i)
```

```
CarName
fueltype
aspiration
doornumber
carbody
drivewheel
enginelocation
enginetype
cylindernumber
fuelsystem
```


```
In [21]: 1 for i in cars.select_dtypes(include='int64'):
2         print(i)
```

```
car_ID
symboling
curbweight
enginesize
horsepower
peakrpm
citympg
highwaympg
```

In [22]:


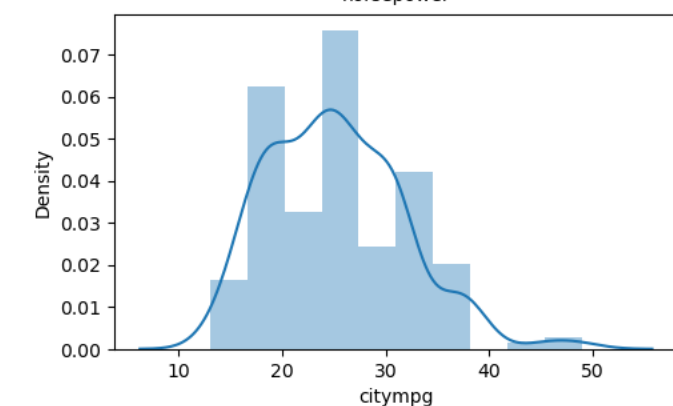

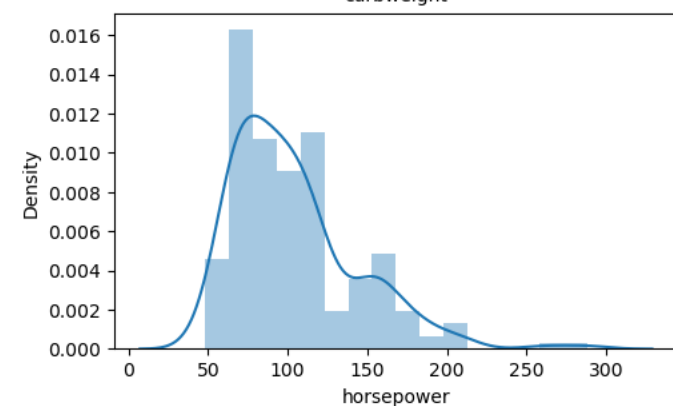




```
1 for i in cars.select_dtypes(include='int64'):
2     print("Column name is:",i)
3     sns.distplot(cars[i])
4     plt.show()
5     print('-----')
```

Column name is: car_ID



In [23]:

```
1 count=1
2 plt.subplots(figsize=(12, 15))
3 for i in cars.select_dtypes(include='int64'):
4     plt.subplot(4,2,count)
5     sns.distplot(cars[i])
6     count+=1
7
8 plt.show()
```



In [24]:

1 cars.head()

Out[24]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	carheight	curbwei
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	21
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	21
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	52.4	21
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	54.3	21
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	54.3	21

In [25]:

1 # numeric dataset
2
3 cars_numeric = cars.select_dtypes(include=['float64', 'int64'])
4 cars_numeric.head()

Out[25]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg
0	1	3	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	9.0	111	5000	21
1	2	3	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	9.0	111	5000	21
2	3	1	94.5	171.2	65.5	52.4	2823	152	2.68	3.47	9.0	154	5000	19
3	4	2	99.8	176.6	66.2	54.3	2337	109	3.19	3.40	10.0	102	5500	24
4	5	2	99.4	176.6	66.4	54.3	2824	136	3.19	3.40	8.0	115	5500	18

In [26]:

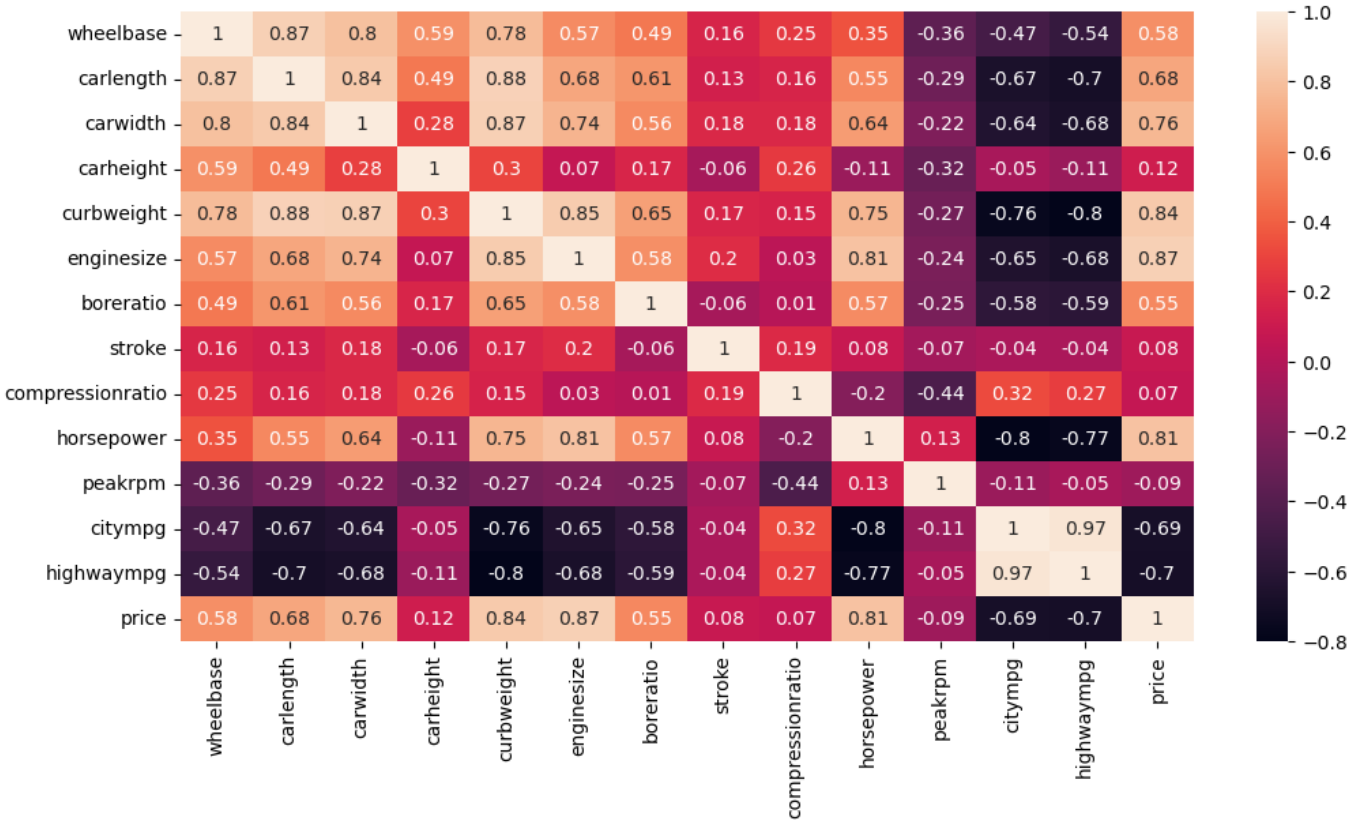
1 cars_numeric = cars_numeric.drop(['car_ID', 'symboling'], axis = 1)
2 cars_numeric.head()

Out[26]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	price
0	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	9.0	111	5000	21	27	13495
1	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	9.0	111	5000	21	27	16500
2	94.5	171.2	65.5	52.4	2823	152	2.68	3.47	9.0	154	5000	19	26	16500
3	99.8	176.6	66.2	54.3	2337	109	3.19	3.40	10.0	102	5500	24	30	13950
4	99.4	176.6	66.4	54.3	2824	136	3.19	3.40	8.0	115	5500	18	22	17450

In [27]:

1 plt.figure(figsize=(12,6))
2 sns.heatmap(round(cars_numeric.corr(), 2), annot = True)
3 plt.show()



```
In [28]: 1 cars.head()
```

```
Out[28]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	carheight	curbwei
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	21
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	21
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	52.4	21
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	54.3	21
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	54.3	21

```
In [29]: 1 cars.CarName.iloc[0]
```

```
Out[29]: 'alfa-romero giulia'
```

```
In [30]: 1 cars.CarName.iloc[0].split(' ')
```

```
Out[30]: ['alfa-romero', 'giulia']
```

```
In [31]: 1 cars.CarName.iloc[3].split(' ')
```

```
Out[31]: ['audi', '100', 'ls']
```

```
In [32]: 1 cars[['Company_Name', 'Model']] = cars['CarName'].str.split(' ', 1, expand = True)
2 cars.head()
```

```
Out[32]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	carheight	curbwei
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	21
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	21
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	52.4	21
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	54.3	21
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	54.3	21

```
In [33]: 1 cars.head()
```

```
Out[33]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	carheight	curbwei
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	21
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	48.8	21
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	52.4	21
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	54.3	21
4	5	2	audi 100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	54.3	21

```
In [34]: 1 cars.insert(3, 'Company_Name', cars.pop('Company_Name'))
2 cars.insert(4, 'Model', cars.pop('Model'))
```

```
In [35]: 1 cars.drop(['CarName'], axis =1, inplace = True)
```

```
In [36]: 1 cars.head()
```

```
Out[36]:
```

	car_ID	symboling	Company_Name	Model	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	c
0	1	3	alfa-romero	giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
1	2	3	alfa-romero	stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
2	3	1	alfa-romero	Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	
3	4	2	audi	100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	
4	5	2	audi	100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	


```
In [37]: 1 cars.Company_Name.unique()
```

```
Out[37]: array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
        'isuzu', 'jaguar', 'maxda', 'mazda', 'buick', 'mercury',
        'mitsubishi', 'Nissan', 'nissan', 'peugeot', 'plymouth', 'porsche',
        'porcshce', 'renault', 'saab', 'subaru', 'toyota', 'toyouta',
        'vokswagen', 'volkswagen', 'vw', 'volvo'], dtype=object)
```

```
In [38]: 1 cars.Company_Name.nunique()
```

```
Out[38]: 28
```

```
In [39]: 1 cars['Company_Name'] = cars['Company_Name'].str.lower()
```

```
In [40]: 1 cars.Company_Name.unique()
```

```
Out[40]: array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
        'isuzu', 'jaguar', 'maxda', 'mazda', 'buick', 'mercury',
        'mitsubishi', 'nissan', 'peugeot', 'plymouth', 'porsche',
        'porcshce', 'renault', 'saab', 'subaru', 'toyota', 'toyouta',
        'vokswagen', 'volkswagen', 'vw', 'volvo'], dtype=object)
```

```
In [41]: 1 cars.Company_Name.nunique()
```

```
Out[41]: 27
```

```
In [42]: 1 def replace_name(a,b):
        2     cars.Company_Name.replace(a,b, inplace=True)
```

```
In [43]: 1 replace_name('maxda', 'mazda')
        2 replace_name('porcshce', 'porsche')
        3 replace_name('toyouta', 'toyota')
        4 replace_name('vokswagen', 'volkswagen')
        5 replace_name('vw', 'volkswagen')
```

```
In [44]: 1 cars.Company_Name.unique()
```

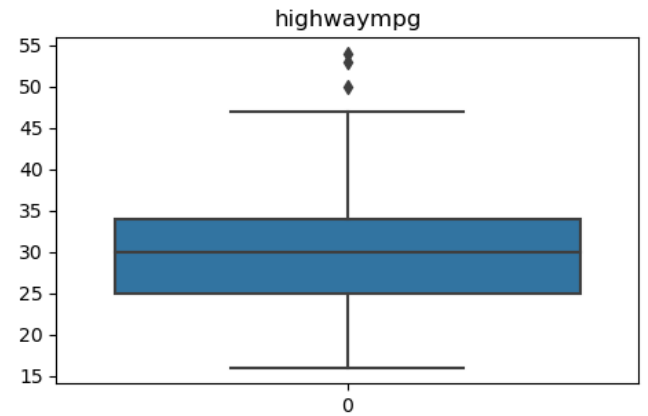
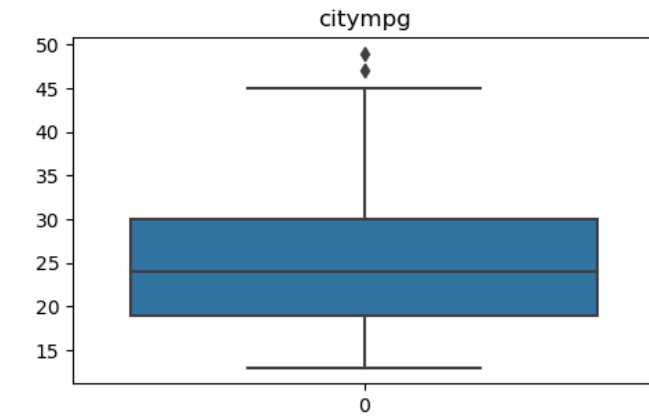
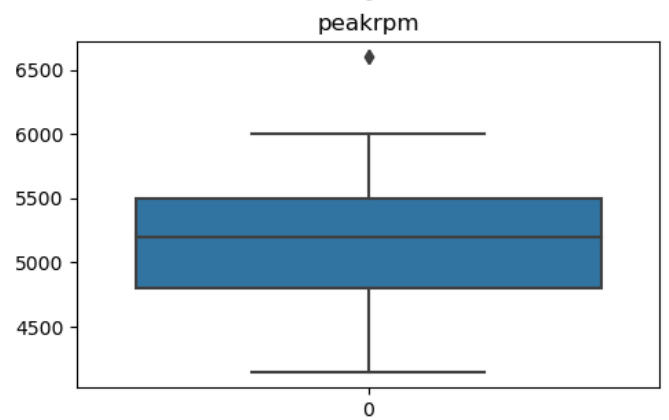


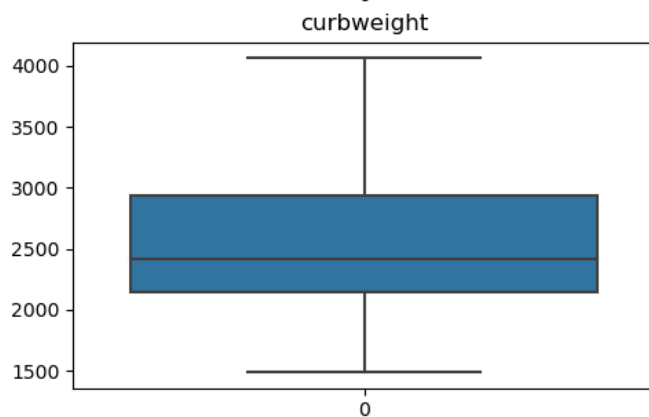

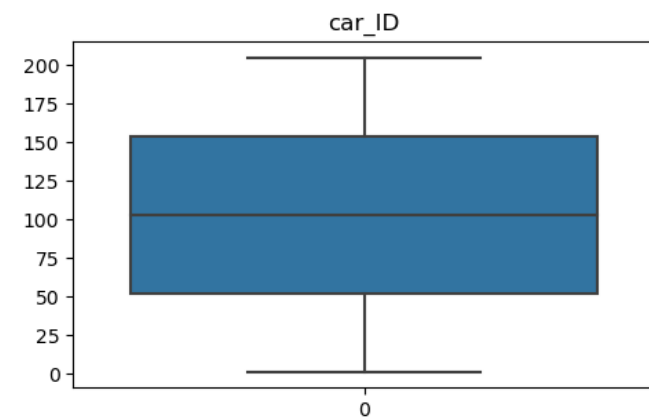
```
Out[44]: array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
        'isuzu', 'jaguar', 'mazda', 'buick', 'mercury', 'mitsubishi',
        'nissan', 'peugeot', 'plymouth', 'porsche', 'renault', 'saab',
        'subaru', 'toyota', 'volkswagen', 'volvo'], dtype=object)
```

```
In [45]: 1 cars.Company_Name.nunique()
```

```
Out[45]: 22
```

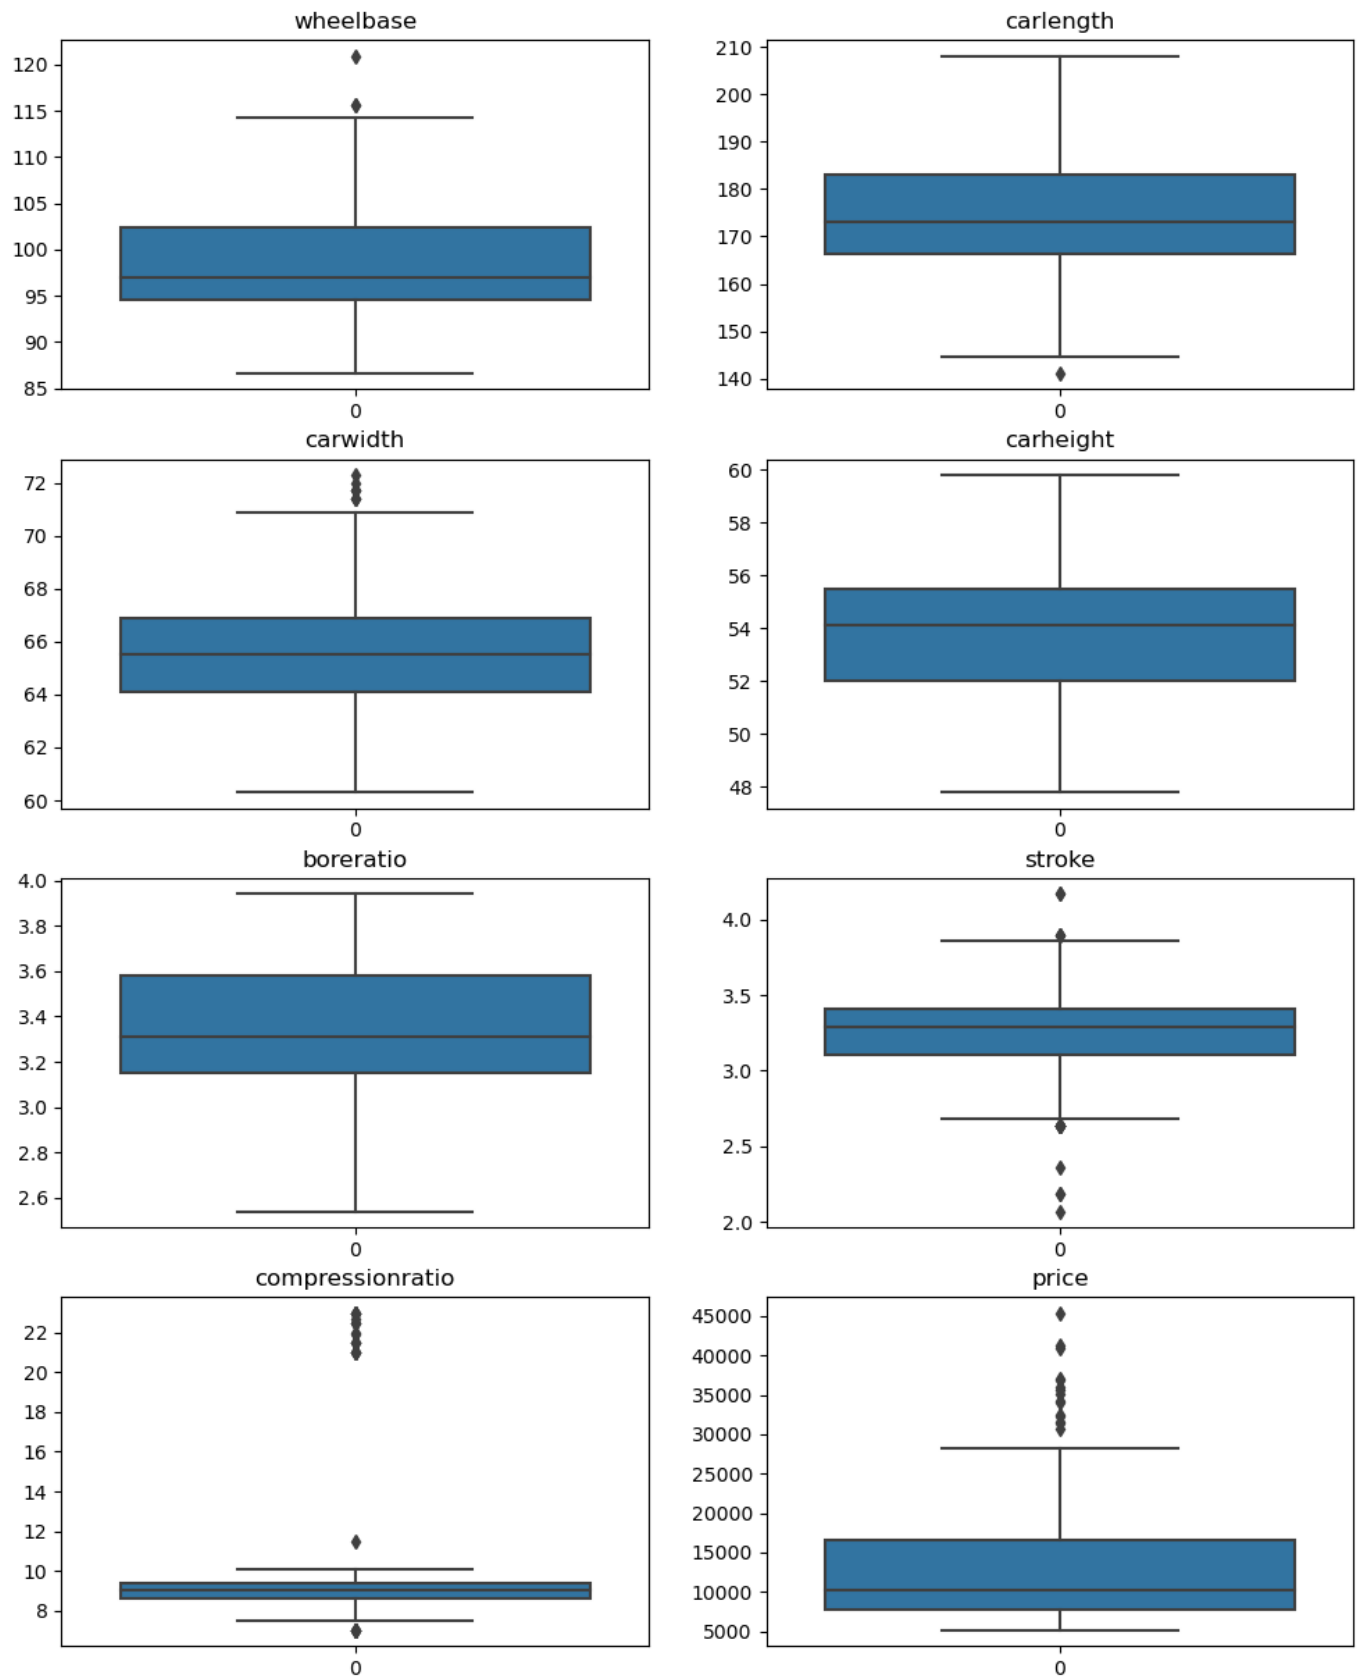
In [46]:

```
1 count=1
2 plt.subplots(figsize=(12, 15))
3 for i in cars.select_dtypes(include='int64'):
4     plt.subplot(4,2,count)
5     sns.boxplot(cars[i])
6     plt.title(i)
7     count+=1
8
9 plt.show()
```



In [47]:

```
1 count=1
2 plt.subplots(figsize=(12, 15))
3 for i in cars.select_dtypes(include='float64'):
4     plt.subplot(4,2,count)
5     sns.boxplot(cars[i])
6     plt.title(i)
7     count+=1
8
9 plt.show()
```



In [48]:


```
1 for i in range(5):
2     print(10**i)
```

```
1
10
100
1000
10000
```


```
In [49]: 1 data = [10**i for i in range(5)]
         2 data
```

```
Out[49]: [1, 10, 100, 1000, 10000]
```

```
In [50]: 1 plt.figure(figsize=(4,3))
         2 plt.plot(data)
         3 plt.show()
```



```
In [51]: 1 plt.figure(figsize=(4,3))
         2 plt.yscale('log')
         3 plt.plot(data)
         4 plt.show()
```



```
In [52]: 1 # map function
         2
         3 my_list = [2,3,4,5,6,7,8,9]
```

```
In [53]: 1 def square(x):
         2     return x*x
```

```
In [54]: 1 list(map(square, my_list))
```

```
Out[54]: [4, 9, 16, 25, 36, 49, 64, 81]
```

```
In [55]: 1 result = list(map(lambda x:x**x, my_list))
         2 result
```

```
Out[55]: [4, 27, 256, 3125, 46656, 823543, 16777216, 387420489]
```

```
In [56]: 1 np.arange(0, 0.00011, 0.00002)
```

```
Out[56]: array([0.e+00, 2.e-05, 4.e-05, 6.e-05, 8.e-05, 1.e-04])
```

```
In [57]: 1 print(format(2.e-05, '.5f'))
```

```
0.00002
```

```
In [58]: 1 range_value = np.arange(0, 0.00011, 0.00002)
```


```
In [59]: 1 def num_conversion(x):
         2     return format(x, '.5f')
```

```
In [60]: 1 print(list(map(num_conversion, range_value)))
```

```
['0.00000', '0.00002', '0.00004', '0.00006', '0.00008', '0.00010']
```

In [61]:

```
1 plt.figure(figsize=(4,3))
2 sns.distplot(cars.price)
3 plt.show()
```



In [62]:

```
1 cars.price.describe(percentiles=[.25, .50, .75, .85, .95, 1])
```

Out[62]:

count	205.000000
mean	13276.710571
std	7988.852332
min	5118.000000
25%	7788.000000
50%	10295.000000
75%	16503.000000
85%	18500.000000
95%	32472.400000
100%	45400.000000
max	45400.000000

Name: price, dtype: float64

In [63]:

```
1 cars.price.median()
```

Out[63]: 10295.0

In [64]:

```
1 cars.price.skew()
```

Out[64]: 1.7776781560914454

In [65]:


```
1 cars.head()
```

Out[65]:

	car_ID	symboling	Company_Name	Model	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	c
0	1	3	alfa-romero	giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
1	2	3	alfa-romero	stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
2	3	1	alfa-romero	Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	
3	4	2	audi	100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	
4	5	2	audi	100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	

In [66]:


```
1 count = 1
2
3 plt.figure(figsize=(12,25))
4 plt.tight_layout()
5 for i in cars.select_dtypes(include=object):
6     plt.subplots_adjust(hspace=0.5)
7     plt.subplot(6,2,count)
8     cars[i].value_counts().plot(kind = 'bar')
9     plt.title(i)
10    count +=1
11
12 plt.show()
```



In [67]:

```
1 sns.boxplot(x = cars.doornumber, y= cars.price)
```

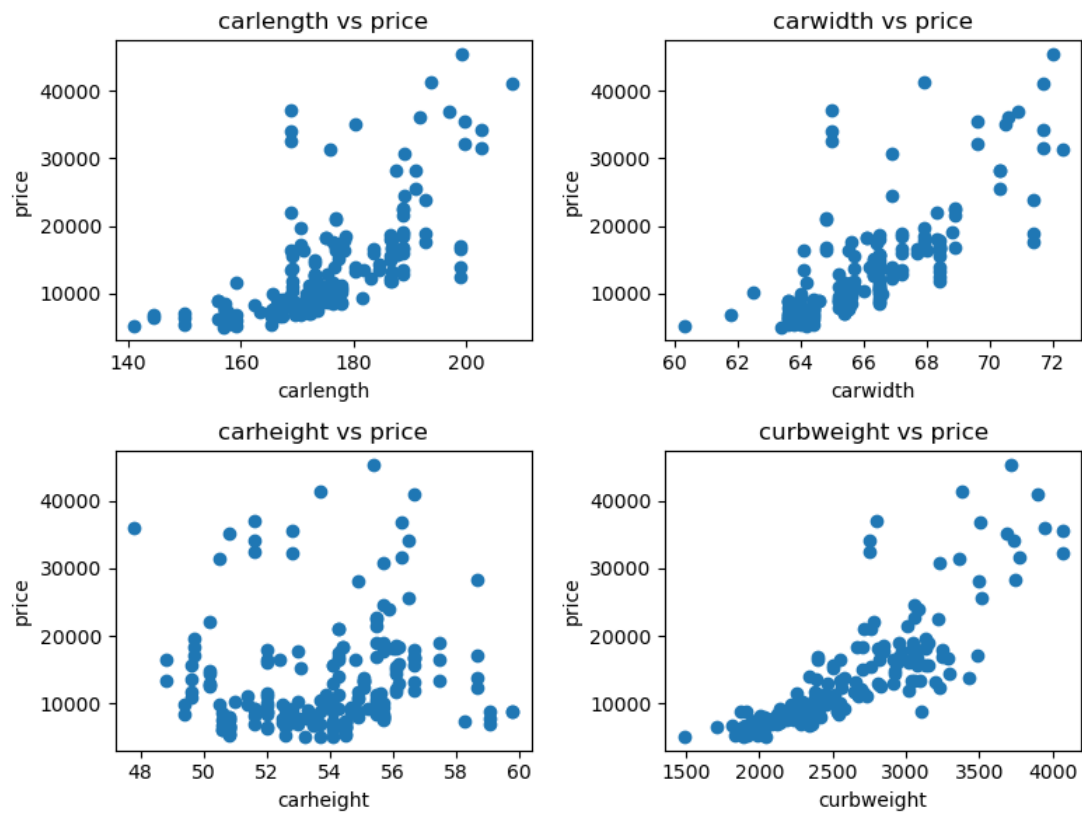
Out[67]: <AxesSubplot:xlabel='doornumber', ylabel='price'>



In [68]:

```
1 #scatter plot
2
3 def scatter_plot(x, fig):
4     plt.subplot(2,2, fig)
5     plt.scatter(cars[x], cars.price)
6     plt.title(x+' vs price')
7     plt.xlabel(x)
8     plt.ylabel('price')
```

```
In [69]: 1 plt.figure(figsize=(8,6))
2
3 scatter_plot('carlength', 1)
4 scatter_plot('carwidth', 2)
5 scatter_plot('carheight', 3)
6 scatter_plot('curbweight', 4)
7
8 plt.tight_layout()
```



Inference :

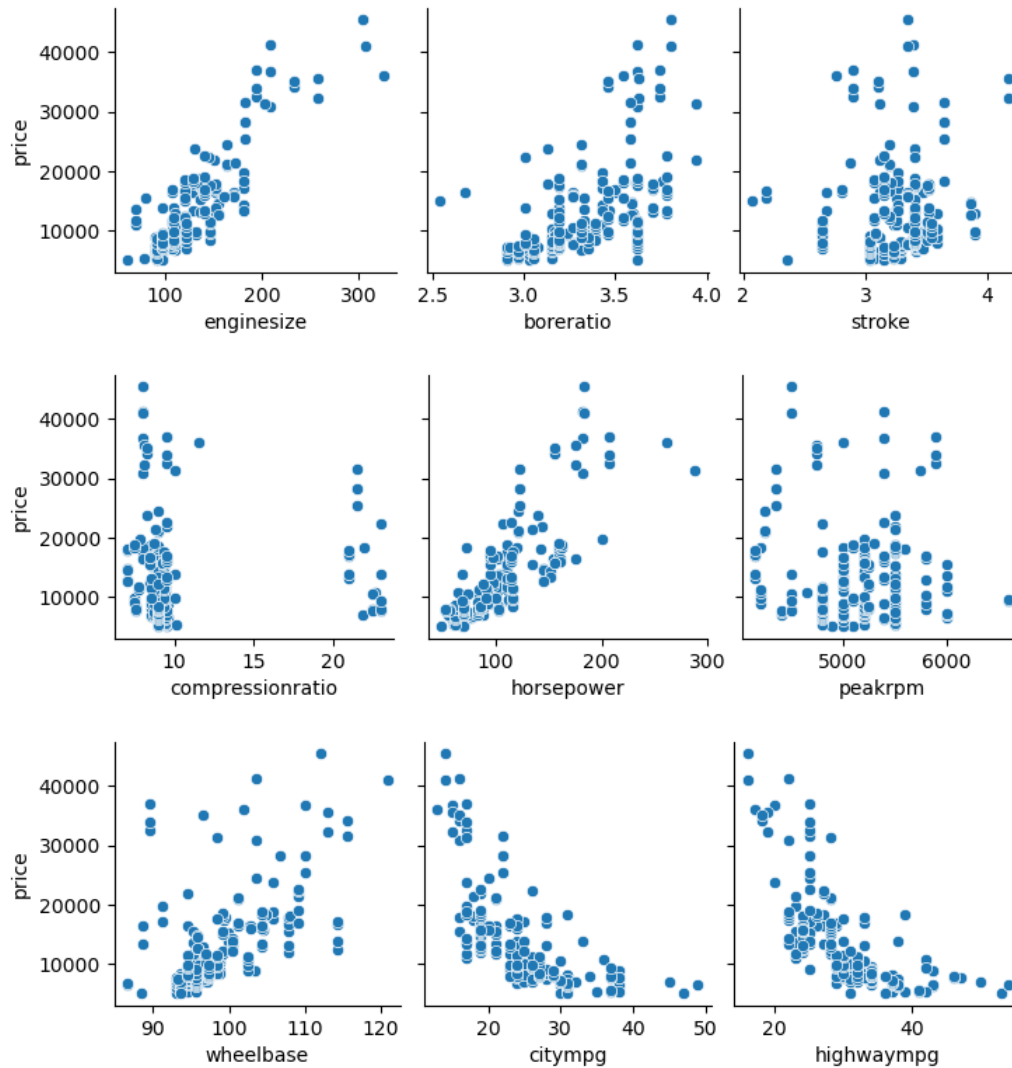
- 1. carwidth, carlength and curbweight seems to have a poitive correlation with price.
- 2. carheight doesn't show any significant trend with price.

```
In [70]: 1 cars.head()
```

Out[70]:

	car_ID	symboling	Company_Name	Model	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	c
0	1	3	alfa-romero	giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
1	2	3	alfa-romero	stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
2	3	1	alfa-romero	Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	
3	4	2	audi	100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	
4	5	2	audi	100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	

```
In [71]: 1 def pair_plot(x, y, z):
2         sns.pairplot(cars, x_vars=[x, y, z], y_vars='price', kind = 'scatter')
3         plt.show()
4
5 pair_plot('engine size', 'bore ratio', 'stroke')
6 pair_plot('compression ratio', 'horsepower', 'peakrpm')
7 pair_plot('wheelbase', 'citympg', 'highwaympg')
```



Inference :

1. engine size, bore ratio, horsepower, wheelbase - seem to have a significant positive correlation with price.
2. citympg, highwaympg - seem to have a significant negative correlation with price.

```
In [72]: 1 cars.price.describe()
```

```
Out[72]: count      205.000000
mean      13276.710571
std       7988.852332
min       5118.000000
25%       7788.000000
50%      10295.000000
75%      16503.000000
max      45400.000000
Name: price, dtype: float64
```

```
In [73]: 1 cars.head()
```

```
Out[73]:
```

	car_ID	symboling	Company_Name	Model	fueltype	aspiration	doornumber	carbody	drivewheel	engine location	wheelbase	carlength	carwidth	c
0	1	3	alfa-romero	giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
1	2	3	alfa-romero	stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
2	3	1	alfa-romero	Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	
3	4	2	audi	100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	
4	5	2	audi	100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	


```
In [74]: 1 temp = cars.copy()
2 table = temp.groupby(['Company_Name'])['price'].mean()
3 table
```

Out[74]: Company_Name
alfa-romero 15498.333333
audi 17859.166714
bmw 26118.750000
buick 33647.000000
chevrolet 6007.000000
dodge 7875.444444
honda 8184.692308
isuzu 8916.500000
jaguar 34600.000000
mazda 10652.882353
mercury 16503.000000
mitsubishi 9239.769231
nissan 10415.666667
peugeot 15489.090909
plymouth 7963.428571
porsche 31400.500000
renault 9595.000000
saab 15223.333333
subaru 8541.250000
toyota 9885.812500
volkswagen 10077.500000
volvo 18063.181818
Name: price, dtype: float64

```
In [75]: 1 temp = temp.merge(table.reset_index(), how = 'left', on='Company_Name')
2 temp.head()
```

Out[75]:

	car_ID	symboling	Company_Name	Model	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	c
0	1	3	alfa-romero	giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
1	2	3	alfa-romero	stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
2	3	1	alfa-romero	Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	
3	4	2	audi	100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	
4	5	2	audi	100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	

```
In [76]: 1 bins = [0, 10000, 20000, 40000]
2 cars_bin = ['budget car', 'avg_priced car', 'highend car']
```

```
In [77]: 1 cars['car_range'] = pd.cut(temp['price_y'], bins, right = False, labels=cars_bin)
2 cars.head()
```

Out[77]:


	car_ID	symboling	Company_Name	Model	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	c
0	1	3	alfa-romero	giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
1	2	3	alfa-romero	stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
2	3	1	alfa-romero	Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	
3	4	2	audi	100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	
4	5	2	audi	100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	

```
In [78]: 1 cars['fuel_economy'] = (0.55*cars['citympg']) + (0.55*cars['highwaympg'])
2 cars.head()
```

Out[78]:

	car_ID	symboling	Company_Name	Model	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	wheelbase	carlength	carwidth	c
0	1	3	alfa-romero	giulia	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
1	2	3	alfa-romero	stelvio	gas	std	two	convertible	rwd	front	88.6	168.8	64.1	
2	3	1	alfa-romero	Quadrifoglio	gas	std	two	hatchback	rwd	front	94.5	171.2	65.5	
3	4	2	audi	100 ls	gas	std	four	sedan	fwd	front	99.8	176.6	66.2	
4	5	2	audi	100ls	gas	std	four	sedan	4wd	front	99.4	176.6	66.4	

```
In [79]: 1 sns.scatterplot(x = cars.fuel_economy, y = cars.price, hue = cars.drivewheel)
2 plt.show()
```



```
In [80]: 1 cars_categorical = cars.select_dtypes(include=object)
2 cars_categorical.head()
```

Out[80]:

	Company_Name	Model	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	enginetype	cylindernumber	fuelsystem
0	alfa-romero	giulia	gas	std	two	convertible	rwd	front	dohc	four	mpfi
1	alfa-romero	stelvio	gas	std	two	convertible	rwd	front	dohc	four	mpfi
2	alfa-romero	Quadrifoglio	gas	std	two	hatchback	rwd	front	ohcv	six	mpfi
3	audi	100 ls	gas	std	four	sedan	fwd	front	ohc	four	mpfi
4	audi	100ls	gas	std	four	sedan	4wd	front	ohc	five	mpfi

```
In [81]: 1 cars_categorical = cars_categorical.drop(['Model'], axis = 1)
2 cars_categorical.head()
```

Out[81]:

	Company_Name	fueltype	aspiration	doornumber	carbody	drivewheel	enginelocation	enginetype	cylindernumber	fuelsystem
0	alfa-romero	gas	std	two	convertible	rwd	front	dohc	four	mpfi
1	alfa-romero	gas	std	two	convertible	rwd	front	dohc	four	mpfi
2	alfa-romero	gas	std	two	hatchback	rwd	front	ohcv	six	mpfi
3	audi	gas	std	four	sedan	fwd	front	ohc	four	mpfi
4	audi	gas	std	four	sedan	4wd	front	ohc	five	mpfi

```
In [82]: 1 cars_categorical = pd.get_dummies(cars_categorical, drop_first= True)
2 cars_categorical.head()
```

Out[82]:

	Company_Name_audi	Company_Name_bmw	Company_Name_buick	Company_Name_chevrolet	Company_Name_dodge	Company_Name_honda	Company_Na
0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0
3	1	0	0	0	0	0	0
4	1	0	0	0	0	0	0

```
In [83]: 1 cars_categorical.shape
```

Out[83]: (205, 50)

```
In [84]: 1 cars_num = cars.select_dtypes(include=['int64', 'float64'])
2 cars_num.head()
```

Out[84]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg
0	1	3	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	9.0	111	5000	21
1	2	3	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	9.0	111	5000	21
2	3	1	94.5	171.2	65.5	52.4	2823	152	2.68	3.47	9.0	154	5000	19
3	4	2	99.8	176.6	66.2	54.3	2337	109	3.19	3.40	10.0	102	5500	24
4	5	2	99.4	176.6	66.4	54.3	2824	136	3.19	3.40	8.0	115	5500	18

```
In [85]: 1 cars = pd.concat([cars_num, cars_categorical], axis = 1)
2 cars.head()
```

Out[85]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg
0	1	3	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	9.0	111	5000	21
1	2	3	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	9.0	111	5000	21
2	3	1	94.5	171.2	65.5	52.4	2823	152	2.68	3.47	9.0	154	5000	19
3	4	2	99.8	176.6	66.2	54.3	2337	109	3.19	3.40	10.0	102	5500	24
4	5	2	99.4	176.6	66.4	54.3	2824	136	3.19	3.40	8.0	115	5500	18

```
In [86]: 1 cars = cars.drop(['car_ID', 'symboling'], axis =1)
2 cars.head()
```

Out[86]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	price
0	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	9.0	111	5000	21	27	13495
1	88.6	168.8	64.1	48.8	2548	130	3.47	2.68	9.0	111	5000	21	27	16500
2	94.5	171.2	65.5	52.4	2823	152	2.68	3.47	9.0	154	5000	19	26	16500
3	99.8	176.6	66.2	54.3	2337	109	3.19	3.40	10.0	102	5500	24	30	13950
4	99.4	176.6	66.4	54.3	2824	136	3.19	3.40	8.0	115	5500	18	22	17450

```
In [87]: 1 cars.shape
```

Out[87]: (205, 65)

```
In [88]: 1 from sklearn.model_selection import train_test_split
2
3 df_train, df_test = train_test_split(cars, train_size=0.7, test_size=0.3, random_state=40)
```

```
In [89]: 1 print(df_train.shape)
2 print(df_test.shape)
```

(143, 65)
(62, 65)

```
In [90]: 1 from sklearn.preprocessing import MinMaxScaler
2
3 scaler = MinMaxScaler()
4 # num_vars = ['wheelbase', 'carlength', 'carwidth', 'carheight', 'curbweight', 'enginesize', 'boreratio', 'stroke',
5 #             'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg', 'price', 'fuel_economy']
6
7 num_vars = ['wheelbase', 'carlength', 'carwidth', 'carheight', 'curbweight', 'enginesize', 'boreratio', 'stroke',
8             'compressionratio', 'horsepower', 'price', 'fuel_economy']
9
10 df_train[num_vars] = scaler.fit_transform(df_train[num_vars])
11 df_train.head()
```

Out[90]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	
145	0.303207	0.431496	0.352941	0.575221	0.338717	0.148438	0.857143	0.271429	0.04375	0.280952	4800	24	29	0
111	0.620991	0.662992	0.647059	0.787611	0.578836	0.195312	0.730159	0.057143	0.08750	0.204762	5000	19	24	0
86	0.282799	0.437795	0.352941	0.336283	0.294093	0.203125	0.642857	0.661905	0.09375	0.171429	5000	25	32	0
113	0.804665	0.855118	0.647059	0.787611	0.668083	0.195312	0.730159	0.057143	0.08750	0.204762	5000	19	24	0
121	0.206997	0.357480	0.196078	0.265487	0.117297	0.078125	0.341270	0.552381	0.15000	0.076190	5500	31	38	0

```
In [91]: 1 X_train = df_train
        2 y_train = df_train.pop('price')
        3 X_train.shape
```

Out[91]: (143, 64)

```
In [92]: 1 X_train.head()
```

Out[92]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	fi
145	0.303207	0.431496	0.352941	0.575221	0.338717	0.148438	0.857143	0.271429	0.04375	0.280952	4800	24	29	
111	0.620991	0.662992	0.647059	0.787611	0.578836	0.195312	0.730159	0.057143	0.08750	0.204762	5000	19	24	
86	0.282799	0.437795	0.352941	0.336283	0.294093	0.203125	0.642857	0.661905	0.09375	0.171429	5000	25	32	
113	0.804665	0.855118	0.647059	0.787611	0.668083	0.195312	0.730159	0.057143	0.08750	0.204762	5000	19	24	
121	0.206997	0.357480	0.196078	0.265487	0.117297	0.078125	0.341270	0.552381	0.15000	0.076190	5500	31	38	

```
In [93]: 1 y_train.head()
```

Out[93]: 145 0.150827
111 0.258301
86 0.074468
113 0.286034
121 0.037234
Name: price, dtype: float64

```
In [94]: 1 num_vars = ['wheelbase', 'carlength', 'carwidth', 'carheight', 'curbweight', 'enginesize', 'boreratio', 'stroke',  
                2 'compressionratio', 'horsepower', 'price', 'fuel_economy']  
        3 df_test[num_vars] = scaler.transform(df_test[num_vars])  
        4 df_test.head()
```

Out[94]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	fi
116	0.620991	0.662992	0.647059	0.787611	0.654059	0.320312	0.920635	0.690476	0.87500	0.204762	4150	28	33	0
5	0.384840	0.514961	0.441176	0.469027	0.337442	0.257812	0.515873	0.633333	0.09375	0.276190	5500	19	25	0
89	0.230321	0.325984	0.196078	0.592920	0.074798	0.105469	0.484127	0.580952	0.15000	0.080952	5200	31	37	0
35	0.288630	0.296063	0.215686	0.592920	0.126222	0.085938	0.293651	0.638095	0.13750	0.114286	6000	30	34	0
185	0.311953	0.426772	0.362745	0.699115	0.212070	0.152344	0.515873	0.633333	0.12500	0.157143	5250	27	34	0

```
In [95]: 1 X_test = df_test
        2 y_test = df_test.pop('price')
```

```
In [96]: 1 X_test.head()
```

Out[96]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	fi
116	0.620991	0.662992	0.647059	0.787611	0.654059	0.320312	0.920635	0.690476	0.87500	0.204762	4150	28	33	
5	0.384840	0.514961	0.441176	0.469027	0.337442	0.257812	0.515873	0.633333	0.09375	0.276190	5500	19	25	
89	0.230321	0.325984	0.196078	0.592920	0.074798	0.105469	0.484127	0.580952	0.15000	0.080952	5200	31	37	
35	0.288630	0.296063	0.215686	0.592920	0.126222	0.085938	0.293651	0.638095	0.13750	0.114286	6000	30	34	
185	0.311953	0.426772	0.362745	0.699115	0.212070	0.152344	0.515873	0.633333	0.12500	0.157143	5250	27	34	

```
In [97]: 1 y_test.head()
```

Out[97]: 116 0.317249
5 0.250093
89 0.007561
35 0.052232
185 0.074618
Name: price, dtype: float64

```
In [98]: 1 print(X_test.shape)
        2 print(y_test.shape)
```

(62, 64)
(62,)

model building

```
In [99]: 1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import mean_squared_error, r2_score
3
4 lm = LinearRegression()
5 lm.fit(X_train, y_train)
```

Out[99]: LinearRegression()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [100]: 1 X_train.head(2)
```

Out[100]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	ft
145	0.303207	0.431496	0.352941	0.575221	0.338717	0.148438	0.857143	0.271429	0.04375	0.280952	4800	24	29	
111	0.620991	0.662992	0.647059	0.787611	0.578836	0.195312	0.730159	0.057143	0.08750	0.204762	5000	19	24	

```
In [101]: 1 y_train.head(2)
```

Out[101]: 145 0.150827
111 0.258301
Name: price, dtype: float64

```
In [102]: 1 X_test.head(2)
```

Out[102]:

	wheelbase	carlength	carwidth	carheight	curbweight	enginesize	boreratio	stroke	compressionratio	horsepower	peakrpm	citympg	highwaympg	ft
116	0.620991	0.662992	0.647059	0.787611	0.654059	0.320312	0.920635	0.690476	0.87500	0.204762	4150	28	33	
5	0.384840	0.514961	0.441176	0.469027	0.337442	0.257812	0.515873	0.633333	0.09375	0.276190	5500	19	25	

```
In [103]: 1 y_test.head(2)
```

Out[103]: 116 0.317249
5 0.250093
Name: price, dtype: float64

```
In [104]: 1 y_train_pred = lm.predict(X_train)
2 y_test_pred  = lm.predict(X_test)
```

```
In [105]: 1 metric = []
2
3 r2_train_lr = r2_score(y_train, y_train_pred)
4 print('r2_train_lr: ', r2_train_lr)
5 metric.append(r2_train_lr)
6
7 r2_test_lr = r2_score(y_test, y_test_pred)
8 print('r2_test_lr: ', r2_test_lr)
9 metric.append(r2_test_lr)
```

r2_train_lr: 0.9697568031056152

r2_test_lr: 0.7834321337362927

Ridge Regression

```
In [106]: 1 from sklearn.linear_model import Ridge
2 from sklearn.model_selection import GridSearchCV
```

```
In [107]: 1 params = {'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1,
2 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
3 4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50, 100, 500, 1000 ]}
4
5 ridge = Ridge()
6
7 folds = 5
8 model_cv = GridSearchCV(estimator= ridge, param_grid=params, scoring='neg_mean_absolute_error', cv = folds,
9 return_train_score= True, verbose = 1)
10
11 model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

```
Out[107]: GridSearchCV(cv=5, estimator=Ridge(),
    param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3,
    0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
    4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50,
    100, 500, 1000]},
    return_train_score=True, scoring='neg_mean_absolute_error',
    verbose=1)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [108]: 1 model_cv.best_params_
```

```
Out[108]: {'alpha': 0.1}
```

```
In [109]: 1 alpha = 0.1
2
3 ridge = Ridge(alpha = alpha)
4 ridge.fit(X_train, y_train)
```

```
Out[109]: Ridge(alpha=0.1)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [110]: 1 y_train_pred = ridge.predict(X_train)
2 y_test_pred = ridge.predict(X_test)
```

```
In [111]: 1 metrics_ridge = []
2
3 r2_score_train_ridge = r2_score(y_train, y_train_pred)
4 print('r2 train ridge: ', r2_score_train_ridge)
5 metrics_ridge.append(r2_score_train_ridge)
6
7 r2_score_test_ridge = r2_score(y_test, y_test_pred)
8 print('r2 train ridge: ', r2_score_test_ridge)
9 metrics_ridge.append(r2_score_test_ridge)
```

```
r2 train ridge:  0.9660447858623624
r2 train ridge:  0.8659775174301652
```

Lasso Regression

```
In [112]: 1 from sklearn.linear_model import Lasso
```

```
In [113]: 1 lasso = Lasso()
2
3 model_cv = GridSearchCV(estimator= lasso, param_grid= params, scoring='neg_mean_absolute_error',
4 cv = folds, return_train_score=True, verbose=1)
5
6 model_cv.fit(X_train, y_train)
```

Fitting 5 folds for each of 28 candidates, totalling 140 fits

```
Out[113]: GridSearchCV(cv=5, estimator=Lasso(),
    param_grid={'alpha': [0.0001, 0.001, 0.01, 0.05, 0.1, 0.2, 0.3,
    0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0, 2.0, 3.0,
    4.0, 5.0, 6.0, 7.0, 8.0, 9.0, 10.0, 20, 50,
    100, 500, 1000]},
    return_train_score=True, scoring='neg_mean_absolute_error',
    verbose=1)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [114]: 1 model_cv.best_params_
```

```
Out[114]: {'alpha': 0.0001}
```

```
In [115]: 1 alpha = 0.0001
          2
          3 lasso = Lasso(alpha = alpha)
          4 lasso.fit(X_train, y_train)
```

Out[115]: Lasso(alpha=0.0001)

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [116]: 1 y_train_pred = lasso.predict(X_train)
          2 y_test_pred = lasso.predict(X_test)
```

```
In [117]: 1 metrics_lasso = []
          2
          3 r2_score_train_lasso = r2_score(y_train, y_train_pred)
          4 print('r2_score_train_lasso: ', r2_score_train_lasso)
          5 metrics_lasso.append(r2_score_train_lasso)
          6
          7 r2_score_test_lasso = r2_score(y_test, y_test_pred)
          8 print('r2_score_test_lasso: ', r2_score_test_lasso)
          9 metrics_lasso.append(r2_score_test_lasso)
```

r2_score_train_lasso: 0.9651911640813117
r2_score_test_lasso: 0.8009345725540462

```
In [127]: 1 # Creating a table which contain all the metrics
          2
          3 table = {'metrics': ['R2 Score (Train)', 'R2 Score (Test)'], 'linear_reg' : metric,
          4         'Ridge_reg': metrics_ridge, 'Lasso_reg':metrics_lasso}
          5 metric_table = pd.DataFrame(table)
          6 metric_table.round(2)
```

Out[127]:

	metrics	linear_reg	Ridge_reg	Lasso_reg
0	R2 Score (Train)	0.97	0.97	0.97
1	R2 Score (Test)	0.78	0.87	0.80

```
In [120]: 1 pd.Series(metrics_ridge, name = 'ridge Regression')
```

Out[120]: 0 0.966045
1 0.865978
Name: ridge Regression, dtype: float64

```
In [ ]: 1
```